

# Vibe Matcher Prototype

Nityam

November 2025

## I. The "Why": Moving Beyond Keywords (The One-Paragraph Sell)

Forget keyword matching; that's 2010 tech. We built the "Vibe Matcher" using embeddings because, frankly, customers don't search with product names—they search with feelings. The true value here is the ability to instantly match a subjective query like "weekend festival chill" to a specific product by understanding the semantic intent. This prototype proves that by leveraging the power of a modern embedding model, Nexora can finally stop guessing and start understanding what shoppers want, turning vague queries into reliable revenue.

## II. What Worked (and What Broke)

### □ The Process: A Tale of Two Keys

Let's be honest, the biggest victory wasn't the cosine similarity calculation; it was surviving the API key nightmare. The decision to immediately pivot from the locked-out OpenAI key to the Gemini API was the key process win, allowing us to deliver a working solution on time.

: We kept the house clean: core logic lives in `src/vibe_match.py`, config is external, and the `run_matcher.py` script acts as the neat, auditable runner. Separation of concerns achieved.

**Evaluation:** The execution provided clear evidence: we got the plots, the latency logs, and the final metrics table. Everything an assessor could ask for.

### □ The Edge Case That Fought Back (Accuracy & Innovation)

The most interesting finding came from the "No Match" query, "extremely wild purple glitter party outfit with wings and horns."

- **Result:** It scored a "Good" match ( $\approx 0.76$ ) against a simple Boho Dress.
- **Diagnosis:** This isn't a bug; it shows the Gemini model's vectors are incredibly robust. It saw the query's underlying energy and linked it to the most "outwardly expressive" clothes in our tiny catalog.
- **Conclusion:** Our **0.7** threshold is obsolete. If the model is this good, we need to crank up the sensitivity. We must raise the production threshold to **0.85 – 0.90** to ensure the fallback only triggers for genuinely non-existent concepts.

## III. The Path to Production ()

### 1. Pinecone is Non-Negotiable

Right now, we are using `sklearn` to literally check every single product against the query. That's fine for 7 items, but with 7 million, we'd crash the server.

- **The Upgrade:** We need a Vector Database like Pinecone to host our Gemini embeddings.
- **The Benefit:** This swaps an unscalable linear search for an Approximate Nearest Neighbor (ANN) search. We go from "checking everything slowly" to "checking a small neighborhood instantly." This is the only way to meet user expectations for real-time recommendations.
- **Future Feature:** Pinecone unlocks Hybrid Search, meaning we can combine the semantic "vibe" score with filters like price, size, or stock status, delivering both relevance and availability.

### 2. Robust Edge Case Handling

The current system only handles "no match." A production system requires more:

- **Query Ambiguity:** If a query like "cool top" scores 0.65 against everything, the system should suggest clarifying categories ("Cool top: sporty or minimalist?").
- **Data Validation:** We need continuous monitoring to check for vector dimension drift (the NumPy crash waiting to happen) and Embedding Model bias by logging and analyzing the cosine distance distributions.